



Telegraph

A peer-to-peer ranking protocol for machine intelligence.

telegraphprotocol.com

Contents.



03	The problem
04	The solution
05	The free flow of intelligence
06	One endpoint
07	Alexandria
08	One terminal, every model
09	Google ranked the web. Bitcoin opened the mine. Telegraph does both, for machines
10	How the ranking works
11	The ranking improves itself
12	It all comes from the people
13	Proof from Season I
14	How we do it

15	The infrastructure
16	What we are not
17	What we are
18	How we attract supply
19	Compete to answer machines
20	The opportunity we create
21	Forecast
22	How we attract demand
23	How we compete
24	Team
25	Raise
26	Thank you

The problem.



The world is going to run on AI in three waves. Humans use it. Then apps plug into it. Then AI transacts with AI, with nobody watching. The intelligence for all three already exists - the demand is arriving, and there is no market between them.

SUPPLY, ALREADY BUILT

Every model ever trained

Every API and toolkit

Every dataset sitting unused

No permissionless market

No truly open ranking

No seamless routing

No easy way in

DEMAND, ARRIVING

Billions of machines

Agents allocating capital

Agents routing logistics

Agents running supply chains


The industry has the intelligence. It has no permissionless way to rank it, route to it, or let people supply it.

The solution.



Telegraph is a peer-to-peer ranking protocol for machine intelligence that supplies answers to humans, apps and machines. We sell one answer, with a receipt: who produced it, what they currently rank for that intent, at what confidence, from \$0.01.


Everything else - the competing miners, the validators, the leaderboard, the consensus, the token - is the factory floor. A manufacturer does not sell its assembly line. We sell the answer that won, and the proof that it won.



The free flow of intelligence.



Requests flow from humans, apps and agents into the network. Miners - people and machines - answer them. Telegraph sits in the middle: it keeps a live ranking of supply, and routes the next request to whoever is winning.




Ranking runs in the background per intent. A paid request is routed from the live leaderboard and returns with a receipt. That call is not scored by 64 validators before it comes back.

One endpoint.



Telegraph gives every human an endless supply of ranked answers, turns every app into an AI company, and makes every machine an expert in everything it wasn't trained on - through one endpoint.




In contrast, \$7.5B OpenRouter only gives an app a model. Telegraph ranks every model for the task and gives the app the answer - and every connected app gets smarter as the network grows.

Alexandria.



OpenAI builds frontier models and puts ChatGPT on top of them. Anthropic does the same with Claude. Telegraph is that structure turned inside out: the protocol does all the AI ranking and its verification, permissionlessly, and serves the result to a terminal we call Alexandria.

The difference is what sits underneath. ChatGPT, Claude and every other model in the world is just a miner on Telegraph, ranked on how accurately it serves the intent behind the query - so the answer comes from whoever is genuinely best at that exact thing, not from whichever vendor you happened to open.




a lab can only serve you what it built - Telegraph serves you whatever won

One terminal, every model.



One query from a human, an app or an agent, routed to whoever is ranked highest for that exact job.



Auto routing

The protocol picks whoever is winning that intent, or select a miner directly.

Every model is supply

ChatGPT, Claude, open-source models, datasets and live APIs all compete as miners.


One front door

Humans, apps and agents query the same endpoint for the same verified answer.

Google ranked the web.
Bitcoin opened the mine.
Telegraph does both, for machines.

Σ

Telegraph is the ranking of intelligence, run as an open competition, completely open by design.




rank the supply, open the competition, and the market builds itself

How the ranking works.



Ranking per intent is the only honest way to capture intelligence from models, datasets, APIs and toolkits, because a provider that is excellent at one task is usually mediocre at the next. Every intent runs its own tournament continuously in the background, so a request never waits on it.




Ranking runs in the background per intent. A paid request is routed from the live leaderboard and returns with a receipt. That call is not scored by 64 validators before it comes back.

Every miner registers with a YAML file that maps its raw JSON response into fixed on-chain slots, so an ordinary Web2 API becomes a deterministic, machine-readable signal that a contract can decode without a parser. That is what makes arbitrary intelligence rankable and consumable by a machine at all.

The ranking improves itself.



The hard problem in ranking is that the ranker can be wrong. Telegraph answers it by making the scoring logic compete as well
- so the standard for being right gets sharper on its own, with nobody deciding.



one Evaluator per intent, replaceable the moment someone builds a better one


the material it is tested on is contributed by anyone, and no contributor can shape more than 2% of it

Evaluators do not start from scratch. They share one trained backbone the whole network reuses, so beating the incumbent is a matter of building a better test - not training a model from nothing.

It all comes from the people.



People integrate the intelligence. Other people build the evaluations for it. A third set verifies the logic ran honestly. Then the machines arrive and consume it. Every part is done by someone who gets paid to do it.



the protocol builds none of it - it opens the door and takes a 2% fee

Our proprietary value is not hidden. A lab's moat is model weights nobody can see; ours is the mechanism itself - the intent registry, permissionless adapters, deterministic scoring, self-promoting evaluation logic and the consensus binding them together. It is fully public and still hard to replicate, the way a well-designed set of rules can be open and still be an edge.

Proof from Season I.



We ran a hackathon, and what we discovered was incredible. All we did was open the doors to a network we promised would handle everything. People came and integrated the intelligence services. Other people - sometimes the same people - built the evaluations for those services. And once the ranking formed, machines came and consumed it all. That is the flywheel, coming from the people, and nothing from Telegraph.

65.2M

Impressions from a roughly \$3K paid push that went viral.

120K

Transactions across the network on testnet.

150

Miners supplying intelligence.

26

Apps live and counting, built during the hackathon.

2,000

Developers registered for Season I.

\$1.80


To convert a developer from X.

Developers who would normally spend three days on oracle plumbing plugged into Telegraph in 30 minutes: one MCP config or Intent ID replaces provider hunting, relayers, fallback logic, escrow contracts and servers they would otherwise maintain forever.

How we do it.




A ranking only means something if nobody can buy their way to the top. A single company ranking intelligence is a company selling placement, so no single company runs the scoring.



The infrastructure.



Miners. The supply side is people and companies, not abstractions.



What we are not.



An L1

We do not run a base chain or compete with Ethereum for settlement. We settle on it.

An L2

A rollup posts its data and proofs to Ethereum so Ethereum can adjudicate its state, which is why rollups are expensive. We post neither, by design, which is why we are free. Security comes from 64 validators and staked Machina.

A data feed or price oracle

An oracle reports a fact that already exists and is observable. We rank judgments that do not exist until someone computes them.

An AI lab

We train nothing, own no models and rent no GPUs. Every model funded anywhere in the world becomes a better supplier at zero cost to us.

A router

A router forwards a request to the cheapest or fastest provider. We decide which is most likely to be right, and prove it afterwards.


An inference reseller

We do not buy capacity and mark it up. We take 2% for ranking, verifying and settling. The 98% never touches our balance sheet.

What we are.



A ranking network for machine intelligence and a settlement rail for humans, apps and machines. The ranking is the only thing we sell, and it is worth nothing the moment a provider can buy their way into it - which is why an open, permissionless network with incentives is the best and most honest way to maintain trust and competition.



neutral by construction - we own no supply, so the ranking is the only thing we sell

Ranking runs in the background per intent. A paid request is routed from the live leaderboard and returns with a receipt. That call is not scored by 64 validators before it comes back.



How we attract supply.

Bitcoin-style money for machines, same fixed supply, pointed at intelligence instead of hashing - rewards go to providers competing to get the answer right. No lab paying fiat could subsidise what Bitcoin achieved through pure token rewards. We do the same thing, but made it a thousand times easier to mine: Bitcoin is niche and expensive, while tens to hundreds of millions of people already own something they could mine Telegraph with today.

	BITCOIN	BITTENSOR (\$TAO)	TELEGRAPH
What you do	Solve cryptographic puzzles.	Run models / provide intelligence in subnets.	Supply any intelligence behind an API (model, data feed, tool, etc.).
Hardware	ASICs (\$4,000 - \$10,000+ per machine).	High-end GPUs or dedicated cloud compute clusters.	Often none. Many miners just use an existing API + a YAML config file.
Ongoing cost	High electricity and cooling (~\$150 - \$600+ per month per machine; ~\$32,000 to \$59,000 in power per BTC mined).	Dynamic subnet slot registration fee (burned in TAO; typically \$200 - \$2,000+) + monthly GPU hosting (\$300 - \$2,000+/month).	Near zero. Main cost is whatever is already paid to run the existing API (\$0 - \$10/month serverless overhead).
Barrier	Very high.	High.	Extremely low.
Paid in	Block rewards + fees (emissions).	Subnet emissions (TAO).	Real demand only. Agents pay USDC, the protocol converts to Machina via TWAP and sends it to the winning miner. Zero emissions.
Total cost to participate	\$4,000 - \$10,000+ upfront per ASIC + \$150 - \$600+/month recurring.	\$500 - \$2,000+ upfront + \$300 - \$2,000+/month recurring compute.	~\$0 upfront + ~\$0 - \$100/month, depending on the kind of intelligence you provide.

Why this matters. Bitcoin and Bittensor both demand capital and continuous spend before you earn anything. Telegraph asks for a config file.

Compete to answer machines.



The same competitive mining structure as Bitcoin, pointed at finding and ranking the best intelligence. Miners compete continuously to answer each intent, validators score those answers with the intent's Evaluator, and an agent's request is routed from the resulting leaderboard to whoever is currently winning.

The token mimics Bitcoin

The token is Machina - Latin for an engine or ingenious contrivance, the word Rome used for the cranes and siege machines that did what people alone could not. 21M hard cap, four-year halvings, zero pre-mine, zero team allocation, zero VC tokens. Miners earn when someone pays. Validators earn emissions for scoring.

But the work has a buyer

Bitcoin burns GPUs on a puzzle that is worthless by design. Miners here compete to produce an answer a machine is programmed to buy. The output is a product instead of heat.

The buyer never touches it

An agent wants an answer, so it pays in stablecoins. The token stays on the miner and validator side, where the competition is - which means every request settled on the network is a forced buy.

The whole world subsidises the supply of intelligence through that competition, and the buyer on Alexandria gets the smallest unit cost of intelligence anywhere.

And paying yourself to farm rewards loses money. The 2% fee on every request is irreversible, and the reward that request can back is capped below it - so recycling your own demand costs more than it returns.

The opportunity we create.



An entirely new layer of the internet for the machine economy - and the first serious attempt at building one.

- A market for measuring intelligence. Evaluation becomes a competitive, replaceable market layer. Anyone who builds a better test can prove it on hidden data and take over scoring for that intent, so the standard for being right keeps getting sharper without anyone owning it.
- A first-mover layer that already works. Nobody has built a ranking, verification and settlement layer for machine intelligence before. This one is running pre-launch, with real requests, real scoring and real settlement.
- Every side of the market gets an upgrade. Model owners get an enormous permissionless market to monetise new and existing AI - no business development, no contracts, no distribution problem. Apps get superpowers without an ML team, a training run or a model decision: a powerful app becomes more powerful and a generic app becomes competitive. And humans and machines stop burning frontier-model tokens on jobs a small specialist does better and cheaper.
- A path to decentralised training. A planned second phase, once the network is live. Structuring the intent and miner data pipelines so cold data becomes valuable signal for model training. For verification integrity, held-out test sets deployed on-chain continuously audit validators, refine miner rankings and drive the false verification rate under 2%.
- Better robots. A robot acts on what it perceives - identify the object, assess the hazard, plan the grasp. Each is an intent. Instead of trusting whatever model shipped inside it, a fleet buys verified answers from whoever is ranked best at that task, and every answer arrives with a receipt it can be audited against later.
- A new machine asset class. Every settled request records which provider answered, what it scored and against what data. That accumulates into an open, auditable ledger of the decisions machines made - provenance that can be priced, audited and traded.

Forecast.



Gartner projects 82 billion AI-influenced transactions per day by 2030. To reach \$500M in annual protocol revenue, Telegraph needs roughly 685 million requests a day - less than 1 in every 119. That is 0.83% of the global agent economy by Year 5.

MODELLED PROTOCOL REVENUE BY YEAR (USD, ANNUAL)



How we attract demand.



Two-sided markets fail the same way: no supply because there is no demand, no demand because there is no supply. We do not wait for either side to arrive.

- Fund the first demand ourselves. The Autonomous Engine acts as Agent Zero, generating continuous real-world requests against a baseline library of standardised intelligence services. Real requests, real scoring, miners paid from real demand. Seed capital with a defined exit.
- Make supply nearly free to join. One config file, no retraining, no migration, competing on the live leaderboard from day one. Season II of the hackathon runs 16 November to 16 December.
- Tell supply exactly what to build. Every request is logged to a public Demand Ledger. Miners do not guess what the market wants. They read what it is already paying for.
- Hardcode the demand rather than advertise for it. Native plugins into LangChain, Eliza and the frameworks agents are already built on put us in the default path.
- Seed one segment at a time. Phase 1 crypto-native. Phase 2 institutional and labs. Phase 3 enterprise network operators. The whole engine inside 12 months.

How we compete.



OpenRouter

A single API into 400+ LLMs, routing on price and latency. Acquired by Stripe for a reported \$7.5B, ~5% fee per call, revenue from \$1M to \$50M in 18 months.

WHERE THEY LEAD TODAY

Live revenue, real developer adoption, brand.

WHY WE WIN

They rank on price, we rank on accuracy per intent - a model can top a general leaderboard and still be mediocre at fraud detection or storm-risk scoring. They deliver raw tokens with no proof; we deliver a scored signal with a receipt an institution can audit. Their supply is curated LLMs; ours is any API-backed service, permissionless. They bill accounts; we settle per request in USDC. And the better a price router gets, the less anyone needs a router.

Chainlink

The decentralised oracle network. A committee observes off-chain data, agrees a median, and attests it on-chain with staking and slashing.

WHERE THEY LEAD TODAY

A decade of integrations plus Swift and DTCC relationships. Distribution we cannot out-argue.

WHY WE WIN

They attest to facts that already exist and are observable, principally prices. We rank judgments that have to be produced - classifications, probabilities, assessments - a surface orders of magnitude larger. Their flagship feeds are largely sponsor-funded, so usage and revenue are only loosely coupled; we meter 2% on every request from the first one. Their operator set is permissioned; ours is permissionless and paid only for winning on accuracy.

Team.



Mark Basa - CEO & Co-founder

Mark leads Telegraph's commercial and business strategy. From 2011 to 2013, Mark built a Bitcoin loyalty app that was accepted to run through two Microsoft incubators. He built a 7-figure agency that operated network nodes. He led all of Bittensor's early PR work, securing 550 earned and organic articles for the protocol, and has advised companies backed by Coinbase Ventures and Uniswap Labs on strategy. He previously served as brand director for a layer-1 blockchain and co-founded a Web3 game studio that raised capital from Japan's largest venture capital firms.

Ahmed Ali - CTO & Co-founder ex- epic

Ahmed is a lead protocol engineer and architect who has spent the past 6 years building and architecting blockchains from scratch. He is the lead engineer behind Telegraph, where he architected the entire protocol - the ranking, evaluation and consensus layers that let any model, API or dataset compete to answer machine requests. Previously, Ahmed designed multiple novel consensus mechanisms and reduced transaction latency from seconds to microseconds. He also built Layer 2 infrastructure by forking the Optimism stack for Epic Chain. Ahmed is a Certified Blockchain Architect with deep expertise in distributed systems and a proven track record of scaling blockchain networks for production.

TELEGRAPH CORE TEAM

Evan Slattery ex- coinbase

Growth

MSc, BEng - aerospace engineering

Haider Ali

Senior Distributed Systems Developer

Backend developer at Basebox 

Umar Gujjar

Web3 & Backend API Engineer

Polygon hackathon winner  - savings protection with Chainlink oracles and Polygon ID

Talha Sajid Chaudhary

Full Stack Developer

MSc Computer Science - built smart contracts for Pear Protocol 

Usman Siddiqui

AI/ML Engineer

MSc Computer Science, silver medal - deep learning, NLP, ML

Muhammad Talha

Integrations Engineer

Raise.



\$10M (\$1.7M committed)

The ask, for 25% equity.


\$30M

Pre-money valuation.

\$40M

Post-money valuation.

Total operating spend across the full 15-month plan - \$10,316,000.



Thank you.



Email

`info@telegraphprotocol.com`

Protocol

`telegraphprotocol.com`

Terminal

`alexandria.telegraphprotocol.com`

